

CSE 6230 Final Project

Optimization of a GPGPU Lattice Boltzmann Method Implementation

Arjun CHINTAPALLI*, Edwin GOH†

Georgia Institute of Technology, Atlanta, GA

CONTENTS

1	Initial Performance Assessment - Grid Size Parameter Sweep	1
2	Initial Performance Assessment - Small Grid Size	2
2.1	BitBucket Commit	2
2.2	Nvprof Results	2
2.3	Profiling with the Nvidia Visual Profiler	2
2.4	Identified Bottlenecks	4
2.4.1	Issue # 1: Kernel is Compute/Memory Latency Bound	4
2.4.2	Issue # 2:	4
2.4.3	Issue # 3: Memory Bandwidth Usage is Limited	6
3	Initial Performance Assessment - Large Grid Size	7
3.1	BitBucket Commit	7
3.2	Nvprof Results	7
3.3	Profiling with the Nvidia Visual Profiler	7
3.4	Identified Bottlenecks	8
3.4.1	Issue # 1: Kernel Performance Is Bound By Instruction And Memory Latency	8
3.4.2	Issue # 2: Low Theoretical Occupancy	8
3.4.3	Issue # 3: Low Memory Bandwidth Utilization	9
4	Notebook Entry # 1: Reducing Memory Accesses with Temporary Local Variables	10
4.1	BitBucket Commit and Modified Code	10
4.2	Proposed Change and Expected Improvements	10
4.3	Profiling Results	11
4.4	Analysis	11
5	Notebook Entry # 2: Changing Grid/Block Size	12
5.1	BitBucket Commit	12
5.2	Proposed Change and Expected Improvements	12
5.3	Profiling Results	13
5.4	Analysis	13
6	Notebook Entry # 3: Removing cudaDeviceSynchronize	14
6.1	BitBucket Commit	14
6.2	Proposed Change and Expected Improvements	14
6.3	Profiling Results	14
6.4	Analysis	15

*School of Computational Science and Engineering

†Graduate Research Assistant, School of Aerospace Engineering

7	Notebook Entry # 4: Push-Pull LBM Implementation	17
7.1	BitBucket Commit	17
7.2	Proposed Change and Expected Improvements	19
7.3	Profiling Results	19
7.4	Analysis	22
8	Final Analysis	24
8.1	Profiling Results from Best-Performing Code	24
8.2	Analysis and Discussion of Improvements and Remaining Bottlenecks	24
8.3	Further Optimization	25
9	Conclusions and Summary	25

1. INITIAL PERFORMANCE ASSESSMENT - GRID SIZE PARAMETER SWEEP

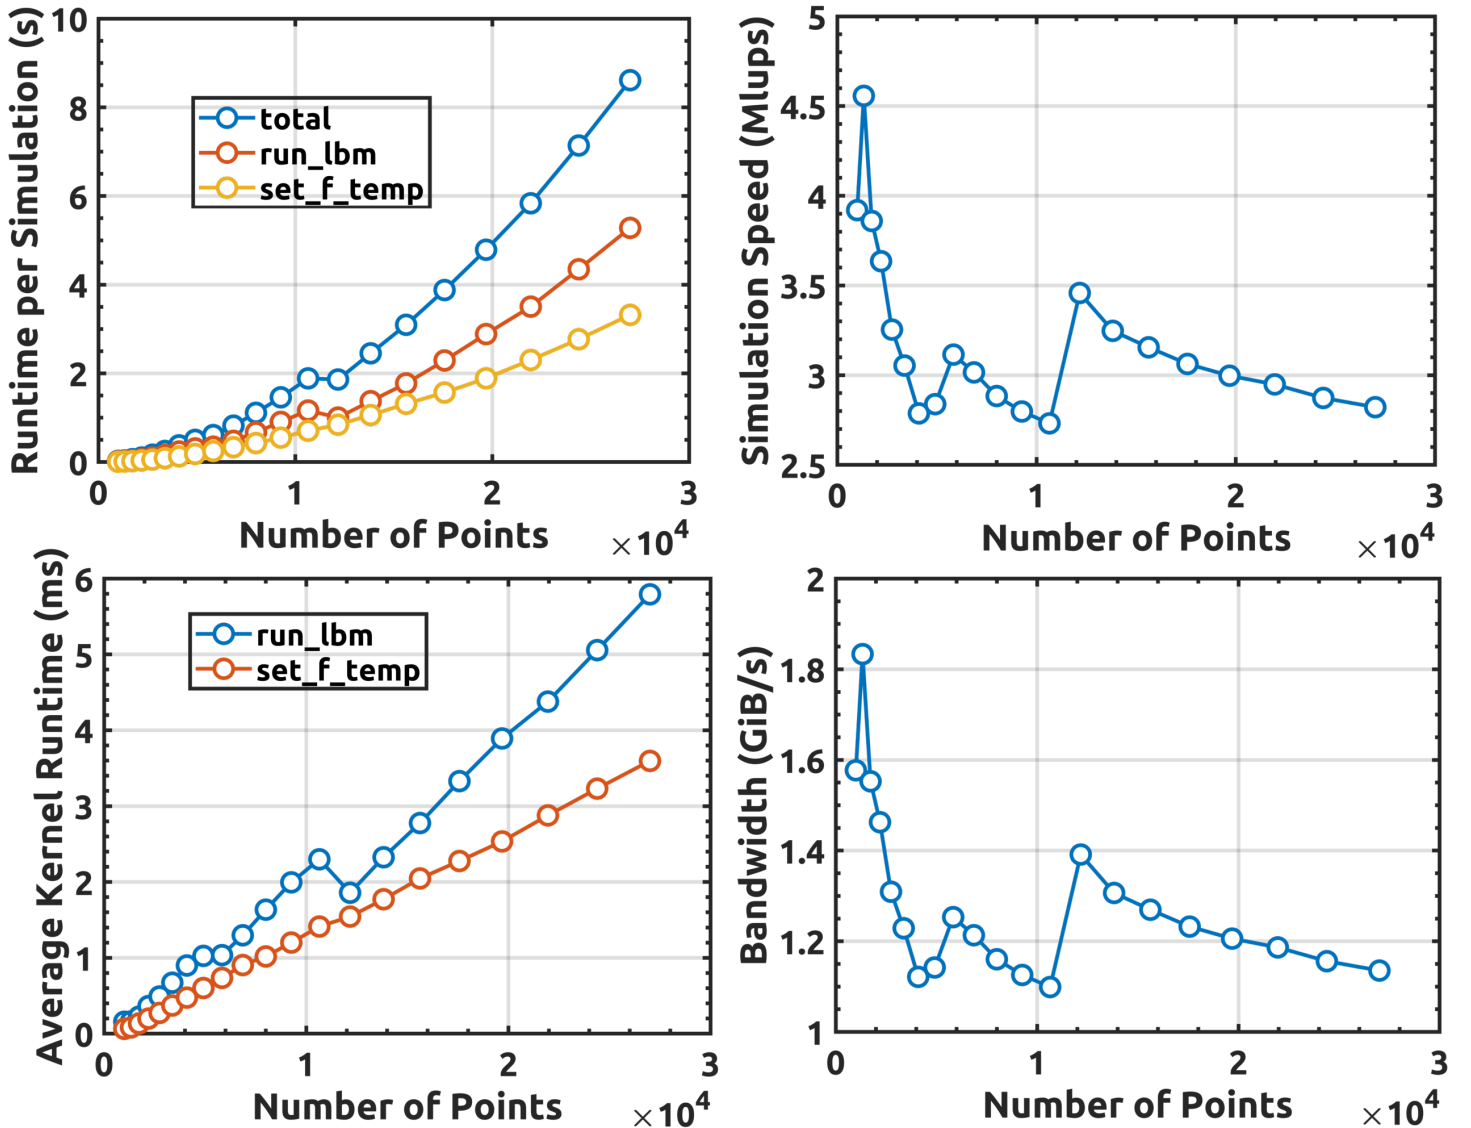


Figure 1: Variation of important performance metrics as a function of number of total points simulated.

Figure 1 shows the results obtained for LBM simulations of different problem sizes. The runtime is obtained by calculating the time taken to complete N^2 iterations of the LBM stream-collide-save algorithm, where N is the number of points in a single spatial dimension. The simulation speed is defined as Mega Lattice-Updates Per Second (Mlups), and is calculated by dividing the total number of points updated (i.e. total number of points \times number of iterations) by the total time taken to run the simulation. Average kernel runtimes are also obtained for the `run_lbm` and `set_f_temp` functions.

The total runtime per simulation increases at a rate that is between quadratic and cubic, i.e. $O(N^{\sim 2.5})$. Under ideal conditions where communication is free and the number of GPU cores is unlimited, the total runtime should only scale as $O(N^2)$ because the number of iterations performed in each iteration is N^2 . In reality, problem sizes that are too large to fit in cache or too large to be executed simultaneously will pose constraints that increase the dependence of the runtime on problem size. In addition, the runtime for the `run_lbm` function scales more strongly than that of `set_f_temp`.

From the average kernel runtime curves, we see that the `set_f_temp` kernel scales almost perfectly linearly with number of points. Although the same amount of work is being done by the kernel, the fact that the average runtime to perform the instructions shown in Listing 1 below increases linear with number of points simulated implies that the time taken to read and write data between `f` and `f_temp`.

Listing 1: Listing of the `set_f_temp` kernel

```

1  __global__ void set_f_temp(PetscReal (*f)[27], double (*f_temp)[27], size_t n) {
2      int x = threadIdx.x;
3      int y = threadIdx.y;
4      int z = blockIdx.x;
5      int pos_index = x + n * (y + n * z); // current (x, y, z)
6      for (int i = 0; i < 27; i++) {
7          f_temp[pos_index][i] = f[pos_index][i];
8      }
9  }
10 }
```

2. INITIAL PERFORMANCE ASSESSMENT - SMALL GRID SIZE

2.1 BitBucket Commit

BitBucket Commit: [d8a7a7ee3077515e94fae2c6da3b738d48e063b0](https://bitbucket.org/isaacgoh/commit/d8a7a7ee3077515e94fae2c6da3b738d48e063b0)

2.2 Nvprof Results

Listing 2: Profiling results obtained using `nvprof` for the small problem size limit.

```

1 [egoh6@ae-comb-305038 parallel] $ nvprof ./test_lbm -num_tests 12 -scale 11
2 ==23667== NVPROF is profiling process 23667, command: ./test_lbm -num_tests 12 -scale 11
3 ==23667== Profiling application: ./test_lbm -num_tests 12 -scale 11
4 ==23667== Profiling result:
5
6   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
7   GPU activities:  64.46%  242.62ms      1452  167.10us  161.06us  181.44us  run_lbm(double[27]*, ...
8   double[27]*, unsigned long, double, IndexBlock)
9   35.24%  132.64ms      1452  91.351us  82.625us  113.57us  set_f_temp(double[27]*, ...
10  double[27]*, unsigned long)
11  0.15%  553.93us        13  42.609us  1.2480us  46.625us  [CUDA memcpy HtoD]
12  0.15%  546.21us        12  45.517us  45.377us  45.632us  [CUDA memcpy DtoH]
13  API calls:  65.95%  569.48ms      1452  392.21us  244.84us  4.5726ms  cudaDeviceSynchronize
14  30.67%  264.85ms        16  16.553ms  18.931us  264.28ms  cudaFree
15  1.79%  15.455ms       2904  5.3210us  4.1840us  35.078us  cudaLaunch
16  0.47%  4.0607ms        27  150.40us  9.7380us  2.5190ms  cudaMalloc
17  0.44%  3.7962ms        12  316.35us  292.55us  342.35us  cudaGetDeviceProperties
18  0.24%  2.0783ms        25  83.130us  14.244us  143.28us  cudaMemcpy
19  0.18%  1.5575ms      11616  134ns      103ns  10.942us  cudaSetupArgument
20  0.09%  798.39us        12  66.532us  61.944us  82.827us  cudaMemGetInfo
21  0.07%  593.41us       2904  204ns      133ns  9.8740us  cudaConfigureCall
22  0.07%  585.77us        185  3.1660us  145ns  130.00us  cuDeviceGetAttribute
23  0.01%  98.325us         2  49.162us  48.774us  49.551us  cuDeviceTotalMem
24  0.01%  91.173us         2  45.586us  36.263us  54.910us  cuDeviceGetName
25  0.00%  12.373us        16  773ns      622ns  1.7790us  cudaEventCreateWithFlags
26  0.00%  11.752us        16  734ns      531ns  2.0820us  cudaEventDestroy
27  0.00%  9.5630us         1  9.5630us  9.5630us  9.5630us  cudaSetDeviceFlags
28  0.00%  5.8580us         2  2.9290us  2.5040us  3.3540us  cudaThreadSynchronize
29  0.00%  4.7290us        11  429ns      285ns  1.1730us  cudaDeviceGetAttribute
30  0.00%  2.0430us         4  510ns      133ns  991ns  cuDeviceGetCount
31  0.00%  1.6080us         1  1.6080us  1.6080us  1.6080us  cudaGetDevice
32  0.00%  1.1700us         3  390ns      234ns  657ns  cuDeviceGet
33  0.00%  662ns          1  662ns      662ns  662ns  cuInit
34  0.00%  508ns          1  508ns      508ns  508ns  cuDriverGetVersion
35  0.00%  459ns          1  459ns      459ns  459ns  cudaGetDeviceCount
```

2.3 Profiling with the Nvidia Visual Profiler

The Nvidia Visual Profiler was used to interactively profile the initial code listed in the above commit.

For preliminary profiling, a local machine (a notebook with a GTX 765M GPU) was used instead of the Bridges cluster in order to decrease turnaround time. Moreover, the Bridges cluster has been found to be inaccessible at times, as shown in the output from `squeue` below:

JOBID	PARTITION	NAME	USER	ST TIME	NODES	NODELIST(REASON)
2057068	GPU-share	lbm-k80	arjunch	PD 0:00	1	(PartitionDown)
2057341	GPU-share	lbm-p100	egoh6	PD 0:00	1	(PartitionDown)

The specifications of this GPU as given by the Nvidia Visual Profiler are shown in Figure 2 below:

[0] GeForce GTX 765M	
GPU UUID	GPU-fe6c21ea-b7dc-9c6c-991c-33e89452d625
Compute Capability	3.0
Max. Threads per Block	1024
Max. Threads per Multiprocessor	2048
Max. Shared Memory per Block	48 KiB
Max. Shared Memory per Multiprocessor	48 KiB
Max. Registers per Block	65536
Max. Registers per Multiprocessor	65536
Max. Grid Dimensions	[2147483647, 65535, 65535]
Max. Block Dimensions	[1024, 1024, 64]
Max. Warps per Multiprocessor	64
Max. Blocks per Multiprocessor	16
Single Precision FLOP/s	1.325 TeraFLOP/s
Double Precision FLOP/s	55.2 GigaFLOP/s
Number of Multiprocessors	4
Multiprocessor Clock Rate	862.5 MHz
Concurrent Kernel	true
Max IPC	7
Threads per Warp	32
Global Memory Bandwidth	64.128 GB/s
Global Memory Size	1.952 GiB
Constant Memory Size	64 KiB
L2 Cache Size	256 KiB
Memcpy Engines	1
PCIe Generation	2
PCIe Link Rate	5 Gbit/s
PCIe Link Width	16

Figure 2: Nvidia Visual Profiler output showing specifications and compute capabilities of the Nvidia GTX 765M

Because the `test_lbm` harness calculates multiple test cases of various grid sizes, only the `run_lbm` kernels which ran smallest grid size (11 x 11 x 11) was analyzed in order to determine the bottlenecks in the limit of small problem sizes. This is shown in Figure 3 below

Analysis Report

run_lbm(double[27]*, double[27]*, unsigned long, double, IndexBlock)

Duration	1.9218 ms (1,921,796 ns)
Grid Size	[11,1,1]
Block Size	[11,11,1]
Registers/Thread	63
Shared Memory/Block	0 B
Shared Memory Requested	48 KiB
Shared Memory Executed	48 KiB
Shared Memory Bank Size	4 B

Figure 3: Nvidia Visual Profiler output showing CUDA kernel being analyzed. The small case (11 x 11 x 11) is being analyzed in this figure.

In this limit of small problem sizes, the initial code uses a block of 121 threads per block, and uses a total of 11 blocks to

update the discretized velocity PDF at each of the $11 \times 11 \times 11$ grid points.

2.4 Identified Bottlenecks

2.4.1 Issue # 1: Kernel is Compute/Memory Latency Bound

Figure 4 shows the utilization of the GPU's compute capabilities and L1 memory during the execution of the `run_lbm` kernel for small problem sizes. The severe underutilization of even a mobile-class laptop GPU indicates that, for small problem sizes, most of the GPU's compute cores L1 memory are idly standing by without receiving any tasks to perform. This may be that there is a latency issue where instructions are either:

1. Taking too long to finish, signifying arithmetic operation latency; or
2. Waiting a long time for the required data to be fetched into or from the GPU, i.e. high memory operation latency.

In other words, the performance of the `run_lbm` kernel when problem sizes are small is limited by the dearth of work to keep the GPU busy.

This further supports our conclusions from Checkpoint 2, where we said our LBM algorithm is bound by memory bandwidth limitations. This issue can be overcome by reducing the need for the kernel to repeatedly access global memory stores, by either prefetching the necessary information or architecting a shared memory system. .




Figure 4: Nvidia Visual Profiler output showing the levels of Compute and L1 Memory utilization for the current limit of small problem sizes.

2.4.2 Issue # 2:

The Nvidia Visual profiler tracks the number of threads, blocks, and warps that are concurrently running during the execution of a kernel, the results of which are shown in Figure 5.

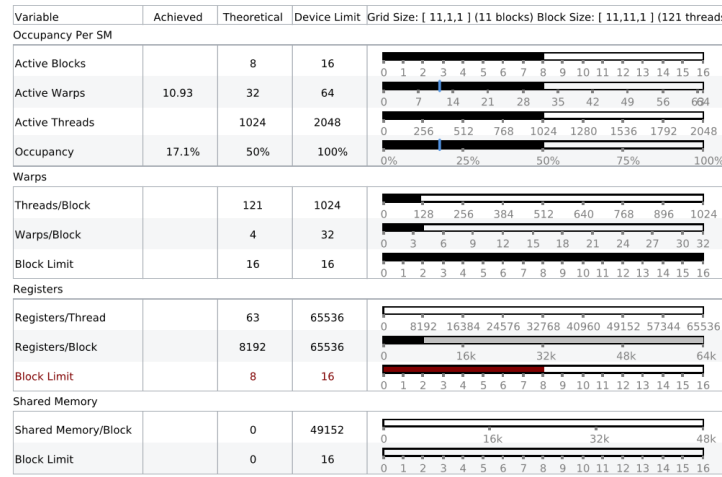


Figure 5: Nvidia Visual Profiler output showing the utilization of blocks, threads, and warps during the execution of the `run_lbm` kernel for small problem sizes.

Based on Figure 5, only 50% of the maximum number of active warps are capable of running concurrently. Even among this theoretical 50%, the current version of the LBM code uses only 10 out of the 32 warps, i.e. 320 threads. For a problem size with $11 \times 11 \times 11 = 1331$ points, the GPU should ideally be able to use 1,331 out of its 2,048 threads to execute the `run_lbm` kernel in one fell swoop. In such a situation the number of active threads would be 1,331, the corresponding number of warps of which is 42. The kernel does not execute enough blocks to hide memory and operation latency. Typically the kernel grid size must be large enough to fill the GPU with multiple "waves" of blocks. Based on theoretical occupancy, device "GeForce GTX 765M" can simultaneously execute 8 blocks on each of the 4 SMs, so the kernel may need to execute a multiple of 32 blocks to hide the compute and memory latency. If the kernel is executing concurrently with other kernels then fewer blocks will be required because the kernel is sharing the SMs with those kernels.

2.4.3 Issue # 3: Memory Bandwidth Usage is Limited

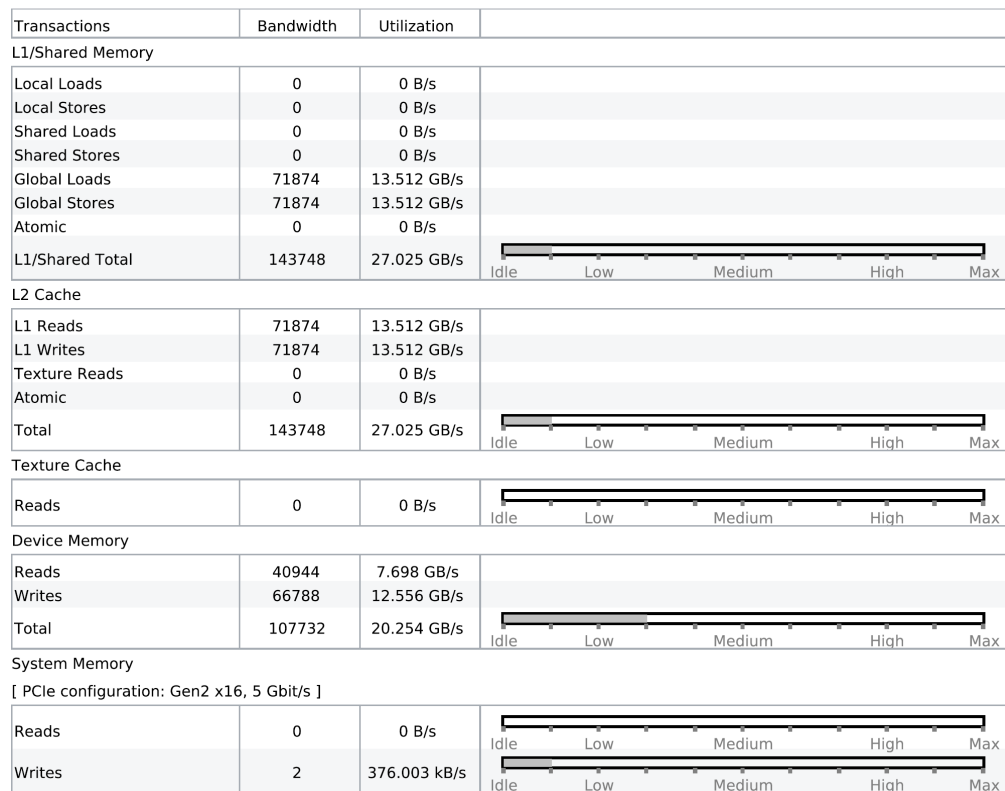


Figure 6: Nvidia Visual Profiler output displaying memory bandwidth utilization for different types of memory available to the GTX 765M device.

Figure 6 obtained from the Nvidia Visual Profiler shows that the `run_lbm` kernel uses only a small amount of the memory bandwidth available on the GTX 765M. Furthermore, there are essentially no operations performed using shared memory. This exclusive reliance on device memory is a limiting factor in terms of performance, and optimization ought to be performed to enable operations using memory shared between threads in a block.

3. INITIAL PERFORMANCE ASSESSMENT - LARGE GRID SIZE

3.1 BitBucket Commit

BitBucket Commit: [d8a7a7ee3077515e94fae2c6da3b738d48e063b0](#)

3.2 Nvprof Results

Listing 3: Profiling results obtained using nvprof for the large problem size limit.

```

1 [egoh6@ae-comb-305038 parallel] $ nvprof ./test_lbm -num_tests 12 -scale 30
2 ==24214== NVPROF is profiling process 24214, command: ./test_lbm -num_tests 12 -scale 30
3 ==24214== Profiling application: ./test_lbm -num_tests 12 -scale 30
4 ==24214== Profiling result:
5
6      Type      Time(%)      Time      Calls      Avg      Min      Max      Name
7      GPU activities:  61.39%  63.3723s    10800  5.8678ms  5.7919ms  6.0652ms  run_lbm(double[27]*, ...,
8                          double[27]*, unsigned long, double, IndexBlock)
9                          38.58%  39.8285s    10800  3.6878ms  3.5961ms  3.8662ms  set_f_temp(double[27]*, ...,
10                             double[27]*, unsigned long)
11
12      Type      Time(%)      Time      Calls      Avg      Min      Max      Name
13      API calls:    99.58%  104.126s    10800  9.6413ms  9.2174ms  14.411ms  cudaDeviceSynchronize
14      0.24%  254.99ms     16  15.937ms  21.038us  252.38ms  cudaFree
15      0.12%  127.30ms    21600  5.8930us  4.1040us  341.67us  cudaLaunch
16      0.03%  26.473ms     25  1.0589ms  14.855us  1.3385ms  cudaMemcpy
17      0.01%  12.040ms    86400   139ns    103ns   304.76us  cudaSetupArgument
18      0.01%  7.8426ms     27  290.47us  9.1790us  2.3155ms  cudaMalloc
19      0.00%  4.5785ms    21600   211ns    153ns   10.775us  cudaConfigureCall
20      0.00%  3.9233ms     12  326.94us  308.64us  366.61us  cudaGetDeviceProperties
21      0.00%  748.70us     12  62.391us  59.361us  75.068us  cudaMemGetInfo
22      0.00%  599.90us    185  3.2420us  132ns   133.07us  cuDeviceGetAttribute
23      0.00%  98.106us      2  49.053us  48.863us  49.243us  cuDeviceTotalMem
24      0.00%  76.516us      2  38.258us  37.724us  38.792us  cuDeviceGetName
25      0.00%  11.156us     16   697ns    518ns   1.8250us  cudaEventCreateWithFlags
26      0.00%  11.074us     16   692ns    457ns   2.1540us  cudaEventDestroy
27      0.00%  9.2610us      1  9.2610us  9.2610us  9.2610us  cudaSetDeviceFlags
28      0.00%  7.4040us      2  3.7020us  3.3660us  4.0380us  cudaThreadSynchronize
29      0.00%  4.8480us     11   440ns    296ns   1.1880us  cudaDeviceGetAttribute
30      0.00%  1.8340us      4   458ns    174ns   1.1890us  cuDeviceGetCount
31      0.00%  1.6130us      1  1.6130us  1.6130us  1.6130us  cudaGetDevice
32      0.00%  1.0260us      3   342ns    166ns    606ns    cuDeviceGet
33      0.00%   556ns      1   556ns    556ns    556ns    cuInit
34      0.00%   475ns      1   475ns    475ns    475ns    cudaGetDeviceCount
35      0.00%   437ns      1   437ns    437ns    437ns    cuDriverGetVersion

```

3.3 Profiling with the Nvidia Visual Profiler

run_lbm(double[27]*, double[27]*, unsigned long, double, IndexBlock)

Duration	5.90161 ms (5,901,607 ns)
Grid Size	[30,1,1]
Block Size	[30,30,1]
Registers/Thread	60
Shared Memory/Block	0 B
Shared Memory Requested	48 KiB
Shared Memory Executed	48 KiB
Shared Memory Bank Size	4 B

Figure 7: Nvidia Visual Profiler output showing the kernel being profiled to be of the problem size 30 x 30 x 30

3.4 Identified Bottlenecks

3.4.1 Issue # 1: Kernel Performance Is Bound By Instruction And Memory Latency




Figure 8: Nvidia Visual Profiler Output showing the levels of Compute and L1 memory utilization for the limit of large problem sizes

3.4.2 Issue # 2: Low Theoretical Occupancy

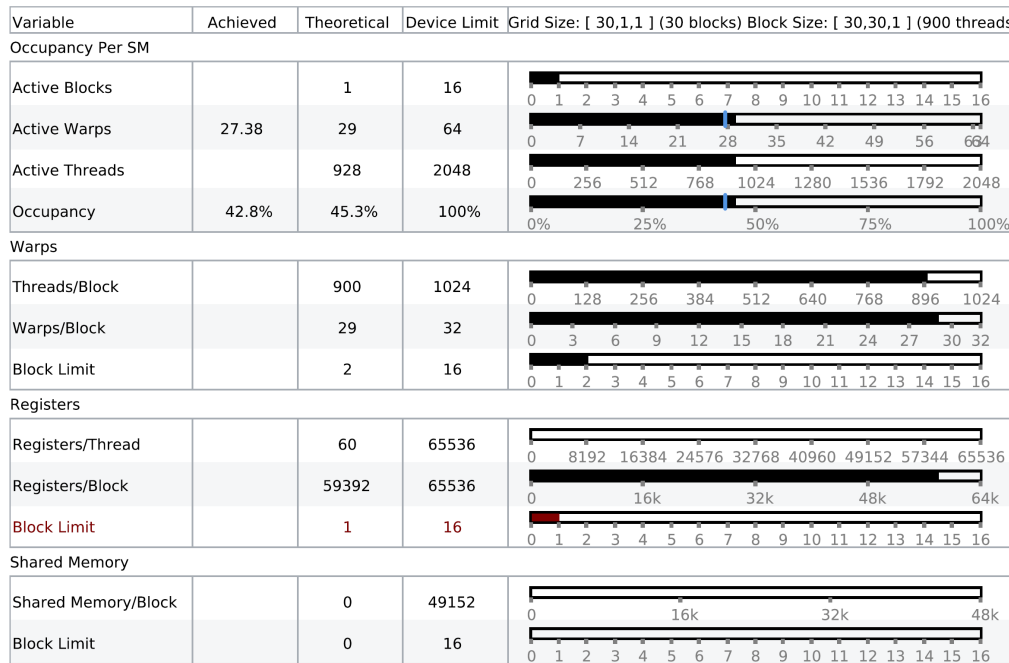


Figure 9: Nvidia Visual Profiler output showing the utilization of blocks, threads, and warps during the execution of the `run_lbm` kernel for small problem sizes.

The large grid size dramatically more so than the smaller grid size demonstrates the memory latency issues inherent with the code. Optimizations with memory latency would speed up the code more so than compute optimizations.

3.4.3 Issue # 3: Low Memory Bandwidth Utilization

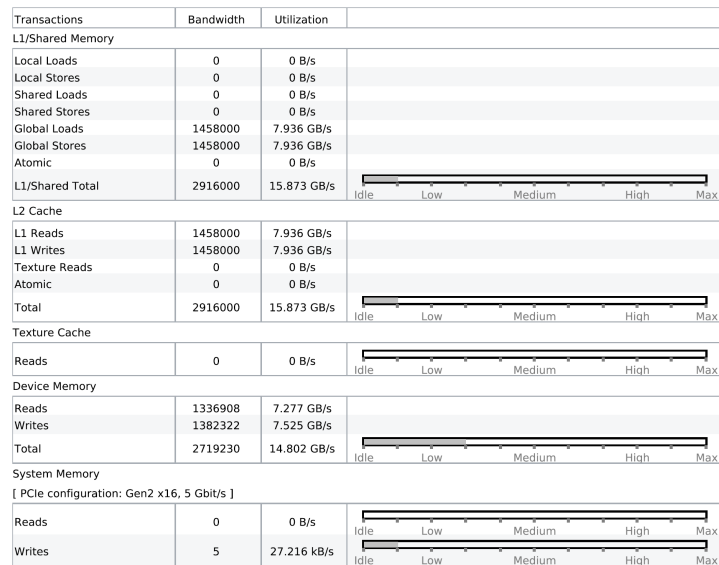


Figure 10: Nvidia Visual Profiler output displaying memory bandwidth utilization for different types of memory available to the GTX 765M device

4. NOTEBOOK ENTRY # 1: REDUCING MEMORY ACCESSES WITH TEMPORARY LOCAL VARIABLES

4.1 BitBucket Commit and Modified Code

BitBucket Commit: [1890a95](#)

```

1  __global__ void run_lbm(PetscReal (*f)[27], double (*f_temp)[27], size_t n, double tau, struct IndexBlock ...
    iBlock)
2  {
3      const double tau_reciprocal = 1/tau;
4      const double one_minus_tau_reciprocal = 1 - tau_reciprocal;
5      int x = threadIdx.x;
6      int y = threadIdx.y;
7      int z = blockIdx.x;
8
9      int pos_index = x + n * (y + n * z); // current (x, y, z) corresponds to a pos_index in the global f array
10
11     double rho = 0, u = 0, v = 0, w = 0;
12     PetscReal f_loc[27];
13
14     for (int dir = 0; dir < 27; dir++) {
15         size_t x_periodic = (n + x - gpu_lbm_velocities[dir][0]) % n;
16         size_t y_periodic = (n + y - gpu_lbm_velocities[dir][1]) % n;
17         size_t z_periodic = (n + z - gpu_lbm_velocities[dir][2]) % n;
18         // Convert Cartesian coordinates to n*n*n row index:
19         size_t old_pos = x_periodic + y_periodic * n + z_periodic * n * n;
20
21         // f[pos_index][dir] = f_temp[old_pos][dir];
22         f_loc[dir]=f_temp[old_pos][dir];
23         rho += f_loc[dir];
24         u += f_loc[dir] * gpu_lbm_velocities[dir][0];
25         v += f_loc[dir] * gpu_lbm_velocities[dir][1];
26         w += f_loc[dir] * gpu_lbm_velocities[dir][2];
27     }
28
29     double rhoinv = 1.0/rho;
30
31     u *= rhoinv; v *= rhoinv; w *= rhoinv;
32     for (int dir = 0; dir < 27; dir++) {
33         // calculate u_alpha (the dot product)
34         double u_alpha = gpu_lbm_velocities[dir][0] * u +
35             gpu_lbm_velocities[dir][1] * v + gpu_lbm_velocities[dir][2] * w;
36         // calculate equilibrium distribution
37         double f_equil_i = gpu_lbm_weights[dir] * rho *
38             (1. + 3. * u_alpha + 4.5 * u_alpha * u_alpha - 1.5*(u*u + v*v + w*w));
39         f[pos_index][dir] = tau_reciprocal * f_equil_i + one_minus_tau_reciprocal * f_loc[dir];
40     }
41 }

```

4.2 Proposed Change and Expected Improvements

Since memory transfers and latencies were determined to be the the primary constraint on our code, optimizations to limit memory read/writes from global device memory were found. The very first initial optimization in this line of thought was to preload each thread's own f distribution values instead of repeatedly accessing them because they don't change throughout the execution of a thread. Priorly, with each iteration a thread would load from global device memory its corresponding lattice nodes f_i and f_{temp}_i distribution values multiple times in the execution, and thus add to the memory dependency latency. An optimization was thus implemented to store these distribution values in a local array with the expected outcome of reducing the time for the `run_lbm` kernel specifically as well as the overall time.

4.3 Profiling Results

Table 1: Raw Performance Statistics

	Optimization 1 (Average across 10 runs)			Base Case (Average across 10 runs)		
N	Total Time (ms)	Avg run_lbm (μ s)	Avg set_ftemp (μ s)	Total Time (ms)	Avg run_lbm (μ s)	Avg set_ftemp (μ s)
10	300.12	194.58	105.54	368.23	266.62	101.61
15	2208.4	508.53	471.92	2436	617.21	465.28
20	8409.731	1074	1027	9865.6	1394.5	1070.1
30	63171.2	3227	3789	84640	5747.7	3924.1

Table 2: Performance Boost Statistics

	Performance Boost		
N	Total Time %	run_lbm %	set_ftemp %
10	0.18	0.27	-0.03
15	0.09	0.17	-0.01
20	0.14	0.23	0.04
30	0.25	0.44	0.03

4.4 Analysis

As expected the profiling results and tables demonstrate that prefetching the f_i distributions dramatically increased performance in the run_lbm kernel specifically as well as the Total Time. Since nothing was modified in the set_ftemp kernel as expected there is no statistically significant increase in performance. In conclusion, as expected reducing the number of device global memory reads lead to performance boosts over 15% at all grid sizes, which further confirms that memory latency is the primary bottleneck.

The logical next step of implementing a shared memory system was also considered but due to the memory limitations of the streaming processors this was deemed infeasible. Consider the base case implementation with the maximum tested grid size of $N = 15$, then each thread block would have 225 threads corresponding to 225 lattice nodes and in total 6075 distributions for each f_i and f_{temp_i} and given the data size of 4 bytes per distribution this would have necessitated a total shared memory of 48,600 bytes but the Nvidia K20 multiprocessors only have 49,152 bytes of shared memory. Additionally with the prefetching implementation, there is only one read and one write per direction of distribution per call to lbm_run, and each thread would read/write information pertaining to its own lattice node, thus shared memory would be of limited utility in this situation.

5. NOTEBOOK ENTRY # 2: CHANGING GRID/BLOCK SIZE

5.1 BitBucket Commit

BitBucket Commit: [b91639b](#)

5.2 Proposed Change and Expected Improvements

In the initial code, the grid-block structure adopted was one where the threads in the block were allocated such that each thread in the x and y dimension represented points in the x and y directions in the simulated grid. Blocks in the CUDA grid's x direction were used to represent points in the z dimension in space. This implementation, however, led to constraints on the number of spatial points in a given dimension, because the maximum number of threads per block is limited to 1,024. Consequently, results such as those shown in Figure 1 only include problem sizes up to 30 x 30 x 30, because larger problem sizes would break the limit on threads per block.

The new grid-block structure can be seen in the following code:

```

1  threads_per_block = PetscMin(prop.maxThreadsPerBlock, threads_per_block);
2  int numBlocks = ceil(N*N*N / threads_per_block + 1);
3  dim3 grid_of_blocks(numBlocks, 1, 1);
4  dim3 block_of_threads(threads_per_block, 1, 1);
5  ...
6  ...
7  int pos_index = threadIdx.x + blockIdx.x * blockDim.x;
8  if (pos_index ≥ n * n * n) {
9      return;
10 }
11 int x = (int) pos_index % (int) n;
12 int y = ((int)pos_index / (int) n) % (int) n;
13 int z = (int) pos_index / ((int) n * (int)n);

```

In the updated version of the LBM code, the `threads_per_block` is a user-defined input to set the number of threads per block. Using this, the user can determine the optimized block size through the `-block_size` input in order to attain maximum performance. The current project found a block size of around 720 to be optimum. This proposed change would increase the occupancy of the code, an issue which was identified in the initial profiling with the Nvidia Visual Profiler. An increased occupancy would result in more threads working simultaneously at any given time, thereby hiding latency and increasing performance.

5.3 Profiling Results

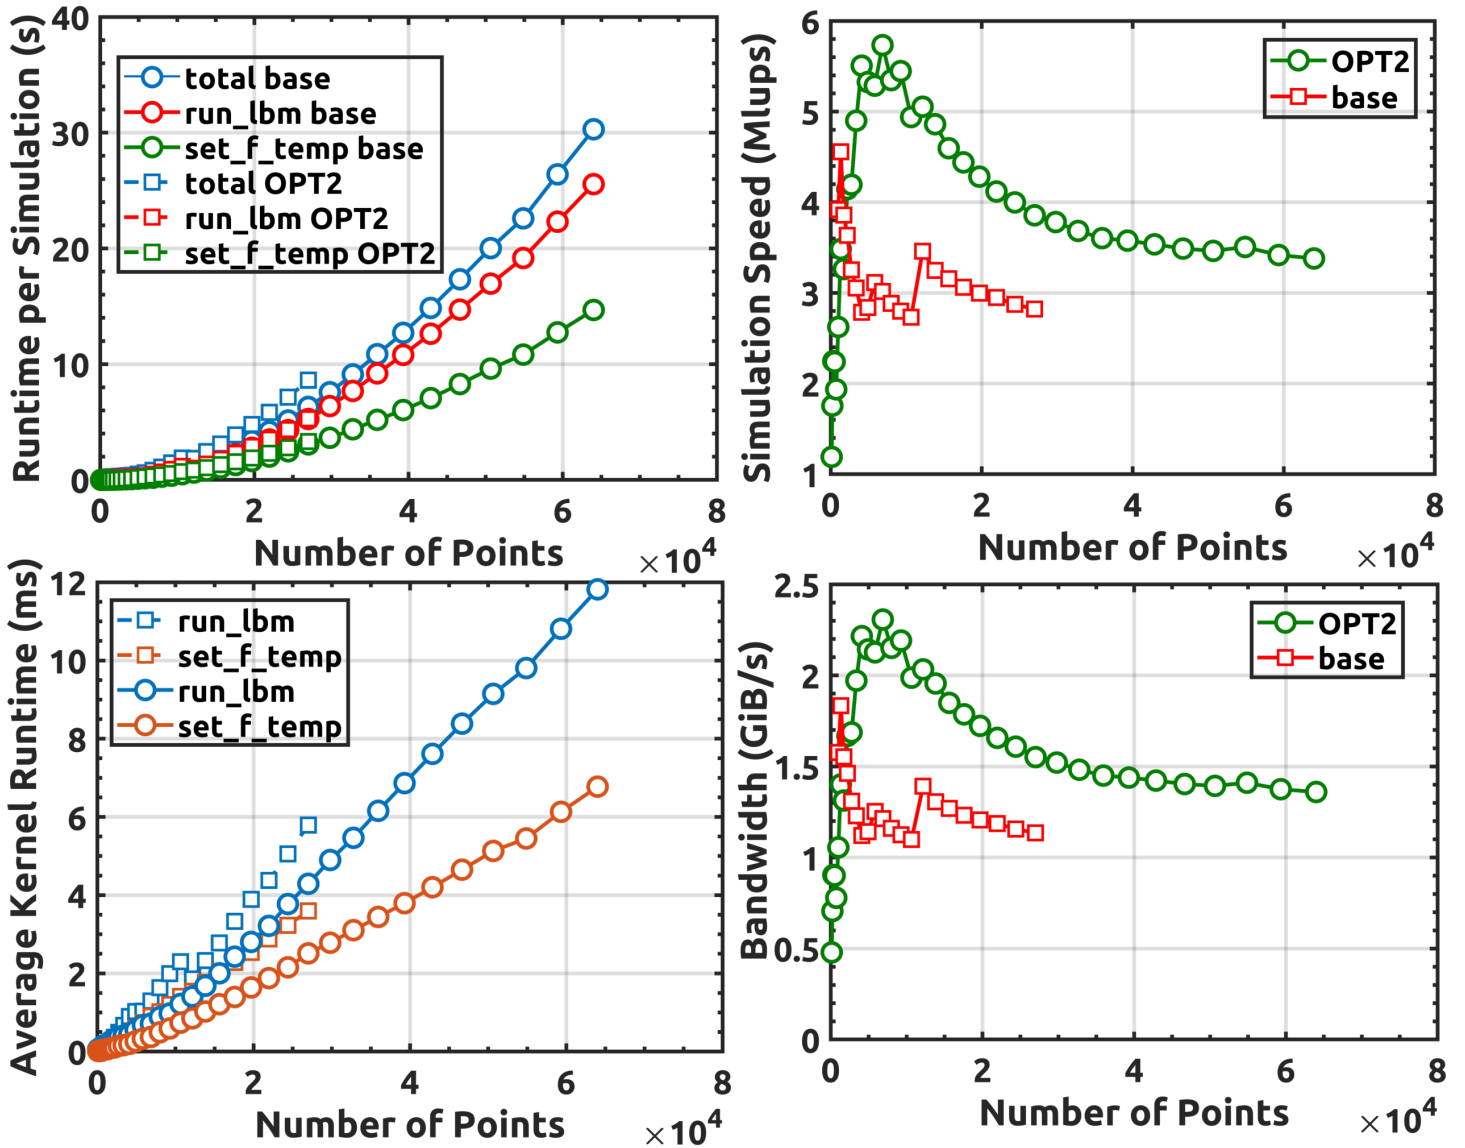


Figure 11: Comparison between the performance of the initial LBM code and the version of the code using the optimized grid-block structure.

5.4 Analysis

Figure 11 above shows the performance from the LBM code using the optimized grid-block structure. It can be seen that not only does the new code, labeled "OPT2" provide decreased runtime, it also increases the range of problem sizes that can be simulated. The increased concurrency results in increased bandwidth, as expected. Furthermore, this change also results in a decreased dependence of average kernel runtimes with the number of points. This might be an artifact resulting from the increased bandwidth due to higher concurrency, which hides the latency behind memory operations in the two kernels.

6. NOTEBOOK ENTRY # 3: REMOVING CUDADEVICESYNCHRONIZE

6.1 BitBucket Commit

BitBucket Commit: [3d0c302](#)

```

1 // LBM_ITERATE_DEVICE MAIN SIMULATION LOOP
2 double begin = seconds();
3 for (int i = 0; i < num_iter; i++) {
4 //   set_f_temp <<< grid_of_blocks, block_of_threads >>> (fLocal[0], f_temp, N); // set f_temp = f //TODO: ...
   move this into run_lbm
5   run_lbm <<<grid_of_blocks, block_of_threads>>> (fLocal[0], f_temp, N, tau, iBlock[0]); // *fLocal[0] is ...
   an array of length 27, with n^3 rows
6 //   cudaDeviceSynchronize();
7 }
```

6.2 Proposed Change and Expected Improvements

With the original code it was noticed that the cudaDeviceSynchronize took up most of the time for the code. Thus it was experimented to see if the function was necessary between iterations.

In a similar train of thought it was also noticed if the set_ftemp function was really necessary and if a work around could be achieved.

6.3 Profiling Results

Listing 4: Profiling results obtained using nvprof for the LBM code without set_f_temp.

```

1 nvprof --log-file "nvprof_no_ftemp_Working.txt" ./test_lbm_noftemp >> "test_lbm_noftemp_Working.txt"
2 Test 7:
3   Test dimensions: [10 x 10 x 10], 100 iterations, tolerance 2.
4   Initializing LBM calculation...
5   Total points = 10 * 10 * 10 = 1000
6   Allocated grid of blocks with dimensions: (10, 1, 1); 10 blocks total
7   Allocated 10 thread blocks with dimensions: (10, 10, 1); 100 threads per block
8   10 blocks * 100 threads per block = 1000 threads total
9 Remote CUDA information:
10 Device name: GeForce GTX 765M
11 # of multiprocessors: 4
12 compute capability: 3.0
13 GPU global memory available: 1999.2 MiB
14
15 Performance Statistics:
16 GPU Memory Allocated: 0.4 (MiB)
17 Runtime: 0.001 (s)
18 Speed: 139.74 (Mlups)
19 Bandwidth: 56.2 (GiB/s)
20   Error 0.128923
21   Passed.
22 Test 8:
23   Test dimensions: [15 x 15 x 15], 225 iterations, tolerance 0.888889
24   Initializing LBM calculation...
25   Total points = 15 * 15 * 15 = 3375
26   Allocated grid of blocks with dimensions: (15, 1, 1); 15 blocks total
27   Allocated 15 thread blocks with dimensions: (15, 15, 1); 225 threads per block
28   15 blocks * 225 threads per block = 3375 threads total
29 Remote CUDA information:
30 Device name: GeForce GTX 765M
31 # of multiprocessors: 4
32 compute capability: 3.0
33 GPU global memory available: 1999.2 MiB
34
35 Performance Statistics:
```



```

36 GPU Memory Allocated: 1.4 (MiB)
37 Runtime: 0.002 (s)
38 Speed: 477.92 (Mlups)
39 Bandwidth: 192.3 (GiB/s)
40     Error 0.0589498
41     Passed.
42 Test 9:
43     Test dimensions: [13 x 13 x 13], 169 iterations, tolerance 1.18343
44     Initializing LBM calculation...
45     Total points = 13 * 13 * 13 = 2197
46     Allocated grid of blocks with dimensions: (13, 1, 1); 13 blocks total
47     Allocated 13 thread blocks with dimensions: (13, 13, 1); 169 threads per block
48     13 blocks * 169 threads per block = 2197 threads total
49 Remote CUDA information:
50 Device name: GeForce GTX 765M
51 # of multiprocessors: 4
52 compute capability: 3.0
53 GPU global memory available: 1999.2 MiB
54
55 Performance Statistics:
56 GPU Memory Allocated: 0.9 (MiB)
57 Runtime: 0.001 (s)
58 Speed: 313.93 (Mlups)
59 Bandwidth: 126.3 (GiB/s)
60     Error 0.904949
61     Passed.
62 ==28035== NVPROF is profiling process 28035, command: ./test_lbm_noftemp
63 ==28035== Profiling application: ./test_lbm_noftemp
64 ==28035== Profiling result:
65      Type      Time(%)      Time      Calls      Avg      Min      Max      Name
66 GPU activities: 99.68% 285.54ms    1121 254.72us 125.73us 615.08us run_lbm(double[27]*, ...
        double[27]*, unsigned long, double, IndexBlock)
67      0.16% 459.97us    11 41.815us 1.3440us 112.42us [CUDA memcpy HtoD]
68      0.16% 445.00us    10 44.499us 18.048us 118.66us [CUDA memcpy DtoH]
69 API calls: 48.53% 320.13ms    14 22.866ms 11.785us 319.63ms cudaFree
70      48.50% 319.92ms    21 15.234ms 15.086us 138.41ms cudaMemcpy
71      1.15% 7.5785ms    23 329.50us 8.5700us 4.9184ms cudaMalloc
72      0.92% 6.0355ms   1121 5.3840us 4.5520us 29.509us cudaLaunch
73      0.50% 3.3096ms    10 330.96us 302.72us 417.78us cudaGetDeviceProperties
74      0.12% 768.36us   5605 137ns    104ns 11.006us cudaSetupArgument
75      0.11% 736.31us   185 3.9800us 143ns 225.30us cuDeviceGetAttribute
76      0.10% 635.25us    10 63.524us 60.417us 71.256us cudaMemGetInfo
77      0.04% 252.68us   1121 225ns    183ns 10.906us cudaConfigureCall
78      0.02% 100.80us     2 50.400us 50.019us 50.782us cuDeviceTotalMem
79      0.01% 74.045us     2 37.022us 33.875us 40.170us cuDeviceGetName
80      0.00% 11.417us    16 713ns    488ns 1.9620us cudaEventDestroy
81      0.00% 11.289us    16 705ns    507ns 2.0260us cudaEventCreateWithFlags
82      0.00% 8.8070us     2 4.4030us 4.3800us 4.4270us cudaThreadSynchronize
83      0.00% 6.7100us     1 6.7100us 6.7100us 6.7100us cudaSetDeviceFlags
84      0.00% 4.5880us    11 417ns    274ns 1.2610us cudaDeviceGetAttribute
85      0.00% 1.7550us     4 438ns    148ns 1.0840us cuDeviceGetCount
86      0.00% 1.7160us     1 1.7160us 1.7160us 1.7160us cudaGetDevice
87      0.00% 1.0450us     3 348ns    180ns 622ns  cuDeviceGet
88      0.00% 603ns        1 603ns    603ns 603ns  cuDriverGetVersion
89      0.00% 550ns        1 550ns    550ns 550ns  cuInit
90      0.00% 448ns        1 448ns    448ns 448ns  cudaGetDeviceCount

```

6.4 Analysis

Although removing the `cudaDeviceSynchronize` was found to dramatically improve the time without affecting the accuracy of the simulations. This is expected because removing the scynchronization reduces the latent time waiting for all blocks to finish. The base code naturally ensures accuracy by keeping the original distributions separate from the distributions post collision.

Removing the `set_ftemp` function was found to work only under certain circumstances. Only if a prior make and run of the code with the `set_ftemp` function included was made, then only would a make and run without the `set_ftemp` function would work. Even with freeing up the prior allocated arrays, an explanation for why the code runs successfully in only some circumstances couldn't be found. Still though the code was found to be faster than the base case, but due to the

unreliability in performance it was not included in the final submission. As expected, the code without `set_ftemp` works better because of the reduction in memory read/writes needed to overwrite the original distributions with each iteration. This reasoning is explained in more detail in the following section.

7. NOTEBOOK ENTRY # 4: PUSH-PULL LBM IMPLEMENTATION

7.1 BitBucket Commit

BitBucket Commit: [228babc](#)

```

1  // MAIN SIMULATION LOOP
2  time_t begin = clock();
3  for (int i = 0; i < num_iter; i++) {
4      push_lbm <<<grid_of_blocks, block_of_threads>>> (fLocal[0], f_temp, N, tau, iBlock[0]); // *fLocal[0] is ...
        an array of length 27, with n^3 rows
5
6      pull_lbm <<<grid_of_blocks, block_of_threads>>> (fLocal[0], f_temp, N, tau, iBlock[0]); // *fLocal[0] is ...
        an array of length 27, with n^3 rows
7  }
8  set_f_temp <<< grid_of_blocks, block_of_threads >>> (f_temp, fLocal[0], N);
9  time_t end = clock();
10 double runtime = (double)(end-begin)/CLOCKS_PER_SEC;
11 double gpu_runtime = 0.001*0.0f;
12
13 size_t doubles_read = 27; // per node every time step
14 size_t doubles_written = 27;
15 size_t doubles_saved = 0; // per node every NSAVE time steps
16
17 // note NX*NY overflows when NX=NY=65536
18 size_t nodes_updated = num_iter*size_t(N*N*N);
19 size_t nodes_saved = 0;
20 double speed = nodes_updated/(1e6*runtime);
21
22 double bandwidth = (nodes_updated*(doubles_read + ...
        doubles_written)+nodes_saved*(doubles_saved))*sizeof(double)/(runtime*bytesPerGiB);
23
24 printf("Performance Statistics: \n");
25 printf("GPU Memory Allocated: %.1f (MiB)\n", 432*N*N*N/bytesPerMiB);
26 printf("Number of Iterations: %u\n", num_iter);
27 printf("Runtime: %.3f (s)\n", runtime);
28 printf("Speed: %.2f (Mlups)\n", speed);
29 printf("Bandwidth: %.1f (GiB/s)\n", bandwidth);
30
31
32
33 PetscFunctionReturn(0);
34 }
35
36 __global__ void set_f_temp(PetscReal (*f)[27], double (*f_temp)[27], size_t n) {
37     int x = threadIdx.x;
38     int y = threadIdx.y;
39     int z = blockIdx.x;
40     int pos_index = x + n * (y + n * z); // current (x, y, z)
41     for (int i = 0; i < 27; i++) {
42         f_temp[pos_index][i] = f[pos_index][i];
43     }
44 }
45
46 __global__ void push_lbm(PetscReal (*f)[27], double (*f_temp)[27], size_t n, double tau, struct IndexBlock ...
        iBlock)
47 {
48     const double tau_reciprocal = 1/tau;
49     const double one_minus_tau_reciprocal = 1 - tau_reciprocal;
50     int x = threadIdx.x;
51     int y = threadIdx.y;
52     int z = blockIdx.x;
53
54     int pos_index = x + n * (y + n * z); // current (x, y, z) corresponds to a pos_index in the global f array
55
56     double rho = 0, u = 0, v = 0, w = 0;
57     PetscReal f_loc[27];

```

```

58 PetscReal f_loc2[27];
59 for (int dir = 0; dir < 27; dir++) {
60     f_loc[dir]=f[pos_index][dir];
61
62     rho += f_loc[dir];
63     u += f_loc[dir] * gpu_lbm_velocities[dir][0];
64     v += f_loc[dir] * gpu_lbm_velocities[dir][1];
65     w += f_loc[dir] * gpu_lbm_velocities[dir][2];
66 }
67
68 double rhoinv = 1.0/rho;
69
70 u *= rhoinv; v *= rhoinv; w *= rhoinv;
71 for (int dir = 0; dir < 27; dir++) {
72
73     size_t x_periodic = (n + x + gpu_lbm_velocities[dir][0]) % n;
74     size_t y_periodic = (n + y + gpu_lbm_velocities[dir][1]) % n;
75     size_t z_periodic = (n + z + gpu_lbm_velocities[dir][2]) % n;
76     // Convert Cartesian coordinates to n*n*n row index:
77     size_t new_pos = x_periodic + y_periodic * n + z_periodic * n * n;
78     // calculate u_alpha (the dot product)
79     double u_alpha = gpu_lbm_velocities[dir][0] * u +
80         gpu_lbm_velocities[dir][1] * v + gpu_lbm_velocities[dir][2] * w;
81     // calculate equilibrium distribution
82     double f_equil_i = gpu_lbm_weights[dir] * rho *
83         (1. + 3. * u_alpha + 4.5 * u_alpha * u_alpha - 1.5*(u*u + v*v + w*w));
84     double temp = tau_reciprocal * f_equil_i + one_minus_tau_reciprocal * f_loc[dir];
85     f_temp[pos_index][dir] = temp;
86     f[new_pos][dir] = temp;
87
88 }
89 }
90 __global__ void pull_lbm(PetscReal (*f)[27], double (*f_temp)[27], size_t n, double tau, struct IndexBlock ...
    iBlock)
91 {
92     const double tau_reciprocal = 1/tau;
93     const double one_minus_tau_reciprocal = 1 - tau_reciprocal;
94     int x = threadIdx.x;
95     int y = threadIdx.y;
96     int z = blockIdx.x;
97
98     int pos_index = x + n * (y + n * z); // current (x, y, z) corresponds to a pos_index in the global f array
99
100     double rho = 0, u = 0, v = 0, w = 0;
101     PetscReal f_loc[27];
102
103     for (int dir = 0; dir < 27; dir++) {
104         size_t x_periodic = (n + x - gpu_lbm_velocities[dir][0]) % n;
105         size_t y_periodic = (n + y - gpu_lbm_velocities[dir][1]) % n;
106         size_t z_periodic = (n + z - gpu_lbm_velocities[dir][2]) % n;
107         // Convert Cartesian coordinates to n*n*n row index:
108         size_t old_pos = x_periodic + y_periodic * n + z_periodic * n * n;
109         f_loc[dir] = f_temp[old_pos][dir];
110         rho += f_loc[dir];
111         u += f_loc[dir] * gpu_lbm_velocities[dir][0];
112         v += f_loc[dir] * gpu_lbm_velocities[dir][1];
113         w += f_loc[dir] * gpu_lbm_velocities[dir][2];
114     }
115
116     double rhoinv = 1.0/rho;
117
118     u *= rhoinv; v *= rhoinv; w *= rhoinv;
119     for (int dir = 0; dir < 27; dir++) {
120         // calculate u_alpha (the dot product)
121         double u_alpha = gpu_lbm_velocities[dir][0] * u +
122             gpu_lbm_velocities[dir][1] * v + gpu_lbm_velocities[dir][2] * w;
123         // calculate equilibrium distribution
124         double f_equil_i = gpu_lbm_weights[dir] * rho *

```

```

125     (1. + 3. * u_alpha + 4.5 * u_alpha * u_alpha - 1.5*(u*u + v*v + w*w));
126     f_temp[pos_index][dir] = tau_reciprocal * f_equl_i + one_minus_tau_reciprocal * f_loc[dir];
127
128 }
129 }

```

7.2 Proposed Change and Expected Improvements

After looking through the literature on implementations of the Lattice Boltzmann Method on GPU's, optimal methods that reduced memory transfers were further considered. Finally a paper by Mawson & Revell¹ was chosen to be implemented because the high performance boost was reported and the implementation didn't seem too ambitious given the time constraints. A limitation that was easily recognized was the need to replace f_temp_i distributions with f_i distributions with every iteration. This replacement is necessary in the base case implementation because a lattice nodes original distributions are required to do the collision and propagation step but the stream step propagates the distribution to neighboring nodes. Thus, 2 distributions need to be maintained to keep track of the original distribution in that iteration and the distribution after collisions. We found an alternate implementation in this article that gets around the need to swap the original and post collision distributions by doing an iteration of pushing the original distributions and an iteration pulling the post-collision distribution. These iterations would avoid the need to have to replace the original distribution array at each iteration. The proposed pseudo code from their work is given for reference.

Algorithm 2 The Push LBM Iteration

```

1: for all Locations  $\mathbf{x}$  in  $f_i(\mathbf{x}, t)$  do
2:   for all  $i$  do
3:     Create a local copy of  $f_i(\mathbf{x}, t)$ 
4:   end for
5:   Apply boundary conditions
6:   Calculate  $\rho$  and  $\mathbf{u}$  using Eqns. [6] and [7].
7:   for all  $i$  do
8:     Calculate  $f^{eq}$  using Eqn. [5]
9:     Perform Collision -  $fLocal_i = fLocal_i(\mathbf{x}, t) + \frac{\Delta t}{\tau} (f^{(eq)}(\mathbf{x}, t) - f(\mathbf{x}, t))$ .
10:    Stream local copies of  $f_i$  to their location  $f_i(\mathbf{x} + \mathbf{e}_i, t + \Delta t)$ 
11:   end for
12: end for

```

Algorithm 3 The Pull LBM Iteration

```

1: for all Locations  $\mathbf{x}$  in  $f_i(\mathbf{x}, t)$  do
2:   for all  $i$  do
3:     Stream to  $fLocal_i$  from the location  $f_i(\mathbf{x} - \mathbf{e}_i \Delta t, t - \Delta t)$ .
4:   end for
5:   Apply boundary conditions.
6:   Calculate  $\rho$  and  $\mathbf{u}$  using Eqns. [6] and [7].
7:   Calculate  $f^{eq}$  using Eqn. [5]
8:   for all  $i$  do
9:     Perform Collision  $fLocal_i = fLocal_i(\mathbf{x}, t) + \frac{\Delta t}{\tau} (f^{(eq)}(\mathbf{x}, t) - f(\mathbf{x}, t))$ .
10:   end for
11: end for

```

Figure 12: Pseudo code describing the algorithm behind the push and pull implementation of the GPU-based LBM implementation

7.3 Profiling Results

Successful implementation of the pseudocode couldn't be achieved as it fails one of the tests in the harness, although the implemented version in the given commit comes close to the tolerance of .88 with an error of .92 for the $n = 15$ gridsize. We

¹Mawson, Mark J., and Alistair J. Revell. "Memory transfer optimization for a lattice Boltzmann solver on Kepler architecture nVidia GPUs." Computer Physics Communications 185.10 (2014): 2566-2574.

believe this implementation is correct except for a small bug which we couldn't figure out. The result from the terminal is given below for validation:

Listing 5: Profiling results obtained using nvprof for the push-pull implementation.

```

1 edwin@Edwin-MacBuntu:~/final/Lattice-Boltzmann-Method/parallel$ make clean && make && nvprof ./test_lbm
2 ==18515== NVPROF is profiling process 18515, command: ./test_lbm
3 Has 1140850688 communicators
4 Running 10 tests of Lattice Boltzmann method with 1 devices
5   Test 0:
6     Test dimensions: [13 x 13 x 13], 169 iterations, tolerance 1.18343
7     Initializing LBM calculation...
8     Total points = 13 * 13 * 13 = 2197
9     Allocated grid of blocks with dimensions: (13, 1, 1); 13 blocks total
10    Allocated 13 thread blocks with dimensions: (13, 13, 1); 169 threads per block
11    13 blocks * 169 threads per block = 2197 threads total
12 Remote CUDA information:
13 Device name: GeForce GT 650M
14 # of multiprocessors: 2
15 compute capability: 3.0
16 GPU global memory available: 976.5 MiB
17
18 Performance Statistics:
19 GPU Memory Allocated: 0.9 (MiB)
20 Number of Iterations: 169
21 Runtime: 0.003 (s)
22 Speed: 110.05 (Mlups)
23 Bandwidth: 44.3 (GiB/s)
24   Error 0.394037
25   Passed.
26   Test 1:
27     Test dimensions: [8 x 8 x 8], 64 iterations, tolerance 3.125
28     Initializing LBM calculation...
29     Total points = 8 * 8 * 8 = 512
30     Allocated grid of blocks with dimensions: (8, 1, 1); 8 blocks total
31     Allocated 8 thread blocks with dimensions: (8, 8, 1); 64 threads per block
32     8 blocks * 64 threads per block = 512 threads total
33 Remote CUDA information:
34 Device name: GeForce GT 650M
35 # of multiprocessors: 2
36 compute capability: 3.0
37 GPU global memory available: 976.5 MiB
38
39 Performance Statistics:
40 GPU Memory Allocated: 0.2 (MiB)
41 Number of Iterations: 64
42 Runtime: 0.002 (s)
43 Speed: 21.73 (Mlups)
44 Bandwidth: 8.7 (GiB/s)
45   Error 0.601047
46   Passed.
47   Test 2:
48     Test dimensions: [8 x 8 x 8], 64 iterations, tolerance 3.125
49     Initializing LBM calculation...
50     Total points = 8 * 8 * 8 = 512
51     Allocated grid of blocks with dimensions: (8, 1, 1); 8 blocks total
52     Allocated 8 thread blocks with dimensions: (8, 8, 1); 64 threads per block
53     8 blocks * 64 threads per block = 512 threads total
54 Remote CUDA information:
55 Device name: GeForce GT 650M
56 # of multiprocessors: 2
57 compute capability: 3.0
58 GPU global memory available: 976.5 MiB
59
60 Performance Statistics:
61 GPU Memory Allocated: 0.2 (MiB)
62 Number of Iterations: 64

```

```
63 Runtime: 0.001 (s)
64 Speed: 24.69 (Mlups)
65 Bandwidth: 9.9 (GiB/s)
66     Error 0.601047
67     Passed.
68 Test 3:
69     Test dimensions: [9 x 9 x 9], 81 iterations, tolerance 2.46914
70     Initializing LBM calculation...
71     Total points = 9 * 9 * 9 = 729
72     Allocated grid of blocks with dimensions: (9, 1, 1); 9 blocks total
73     Allocated 9 thread blocks with dimensions: (9, 9, 1); 81 threads per block
74     9 blocks * 81 threads per block = 729 threads total
75 Remote CUDA information:
76 Device name: GeForce GT 650M
77 # of multiprocessors: 2
78 compute capability: 3.0
79 GPU global memory available: 976.5 MiB
80
81 Performance Statistics:
82 GPU Memory Allocated: 0.3 (MiB)
83 Number of Iterations: 81
84 Runtime: 0.003 (s)
85 Speed: 19.10 (Mlups)
86 Bandwidth: 7.7 (GiB/s)
87     Error 0.497029
88     Passed.
89
90 Test 5:
91     Test dimensions: [8 x 8 x 8], 64 iterations, tolerance 3.125
92     Initializing LBM calculation...
93     Total points = 8 * 8 * 8 = 512
94     Allocated grid of blocks with dimensions: (8, 1, 1); 8 blocks total
95     Allocated 8 thread blocks with dimensions: (8, 8, 1); 64 threads per block
96     8 blocks * 64 threads per block = 512 threads total
97 Remote CUDA information:
98 Device name: GeForce GT 650M
99 # of multiprocessors: 2
100 compute capability: 3.0
101 GPU global memory available: 976.5 MiB
102
103 Performance Statistics:
104 GPU Memory Allocated: 0.2 (MiB)
105 Number of Iterations: 64
106 Runtime: 0.001 (s)
107 Speed: 24.27 (Mlups)
108 Bandwidth: 9.8 (GiB/s)
109     Error 0.601047
110     Passed.
111 Test 6:
112     Test dimensions: [8 x 8 x 8], 64 iterations, tolerance 3.125
113     Initializing LBM calculation...
114     Total points = 8 * 8 * 8 = 512
115     Allocated grid of blocks with dimensions: (8, 1, 1); 8 blocks total
116     Allocated 8 thread blocks with dimensions: (8, 8, 1); 64 threads per block
117     8 blocks * 64 threads per block = 512 threads total
118 Remote CUDA information:
119 Device name: GeForce GT 650M
120 # of multiprocessors: 2
121 compute capability: 3.0
122 GPU global memory available: 976.5 MiB
123
124 Performance Statistics:
125 GPU Memory Allocated: 0.2 (MiB)
126 Number of Iterations: 64
127 Runtime: 0.002 (s)
128 Speed: 20.24 (Mlups)
129 Bandwidth: 8.1 (GiB/s)
130     Error 0.601047
```

```

131     Passed.
132 Test 7:
133     Test dimensions: [10 x 10 x 10], 100 iterations, tolerance 2.
134     Initializing LBM calculation...
135     Total points = 10 * 10 * 10 = 1000
136     Allocated grid of blocks with dimensions: (10, 1, 1); 10 blocks total
137     Allocated 10 thread blocks with dimensions: (10, 10, 1); 100 threads per block
138     10 blocks * 100 threads per block = 1000 threads total
139 Remote CUDA information:
140 Device name: GeForce GT 650M
141 # of multiprocessors: 2
142 compute capability: 3.0
143 GPU global memory available: 976.5 MiB
144
145 Performance Statistics:
146 GPU Memory Allocated: 0.4 (MiB)
147 Number of Iterations: 100
148 Runtime: 0.002 (s)
149 Speed: 49.85 (Mlups)
150 Bandwidth: 20.1 (GiB/s)
151     Error 0.415925
152     Passed.
153 Test 8:
154     Test dimensions: [15 x 15 x 15], 225 iterations, tolerance 0.888889
155     Initializing LBM calculation...
156     Total points = 15 * 15 * 15 = 3375
157     Allocated grid of blocks with dimensions: (15, 1, 1); 15 blocks total
158     Allocated 15 thread blocks with dimensions: (15, 15, 1); 225 threads per block
159     15 blocks * 225 threads per block = 3375 threads total
160 Remote CUDA information:
161 Device name: GeForce GT 650M
162 # of multiprocessors: 2
163 compute capability: 3.0
164 GPU global memory available: 976.5 MiB
165
166 Performance Statistics:
167 GPU Memory Allocated: 1.4 (MiB)
168 Number of Iterations: 225
169 Runtime: 0.005 (s)
170 Speed: 167.45 (Mlups)
171 Bandwidth: 67.4 (GiB/s)
172     Error 0.92373
173 [0]PETSC ERROR: ----- Error Message ...
174 [0]PETSC ERROR: Error in external library
175 [0]PETSC ERROR: Test 8 failed residual test at threshold 0.888889 with value 0.92373
176
177 system msg for write_line failure : Bad file descriptor
178 ==18515== Profiling application: ./test_lbm
179 ==18515== Profiling result:
180 Time(%)      Time      Calls      Avg      Min      Max      Name
181 50.61%  339.08ms      952  356.18us  165.72us  713.42us  pull_lbm(double[27]*, double[27]*, unsigned long, ...
182         double, IndexBlock)
183 49.07%  328.82ms      952  345.40us  162.05us  675.25us  push_lbm(double[27]*, double[27]*, unsigned long, ...
184         double, IndexBlock)
185 0.20%   1.3705ms         9  152.28us  76.734us  461.53us  set_f_temp(double[27]*, double[27]*, unsigned long)
186 0.06%   391.19us        10  39.119us  1.5680us  115.90us  [CUDA memcpy HtoD]
187 0.06%   373.53us         9  41.503us  18.559us  114.81us  [CUDA memcpy DtoH]
188
189 ===== Error: Application returned non-zero code 76

```

7.4 Analysis

Although this implementation fails for gridsizes larger than 15, the below results were achieved for a gridsize of 10. The total time for this implementation was 293.7 ms in comparison to a total time of 368.23 ms for the base case code on the same machine. Thus, this implementation would have led to a speed up in the performance in comparison to the base case.

This optimization though performs better than the first optimization with prefetching combined with the base case, which achieved a total time of 300.12 ms. Thus, if the bug contributing to elevated errors could have been figured out the push pull implementation would have been included in the final.

We believe our implementation of the given psuedocode fails at larger test sizes due to a mixup between the f_i and f_{temp_i} distributions somewhere in the code. Even as is though, at smaller test sizes this code is better than the first optimization implemented and holds good potential if debugged. The improvement in performance that can already be seen is due to the removal of the need to read and write over an array the size of the number of points with each iteration.

Listing 6: Results for gridsize 10 obtained using nvprof for the push-pull implementation.

```

1  ==19702== Profiling application: ./test_lbm -scale 10
2  ==19702== Profiling result:
3  Time(%)      Time       Calls     Avg       Min       Max    Name
4  60.02%   176.58ms       1000   176.58us  172.70us  223.68us  push_lbm(double[27]*, double[27]*, unsigned long, ...
        double, IndexBlock)
5  39.45%   116.07ms       1000   116.07us  112.96us  125.18us  pull_lbm(double[27]*, double[27]*, unsigned long, ...
        double, IndexBlock)
6   0.27%    804.78us         10   80.478us  79.935us  80.990us  set_f_temp(double[27]*, double[27]*, unsigned long)
7   0.13%    385.15us         11   35.013us  1.4400us  45.055us  [CUDA memcpy HtoD]
8   0.12%    356.12us         10   35.612us  34.943us  36.608us  [CUDA memcpy DtoH]

```

8. FINAL ANALYSIS

8.1 Profiling Results from Best-Performing Code

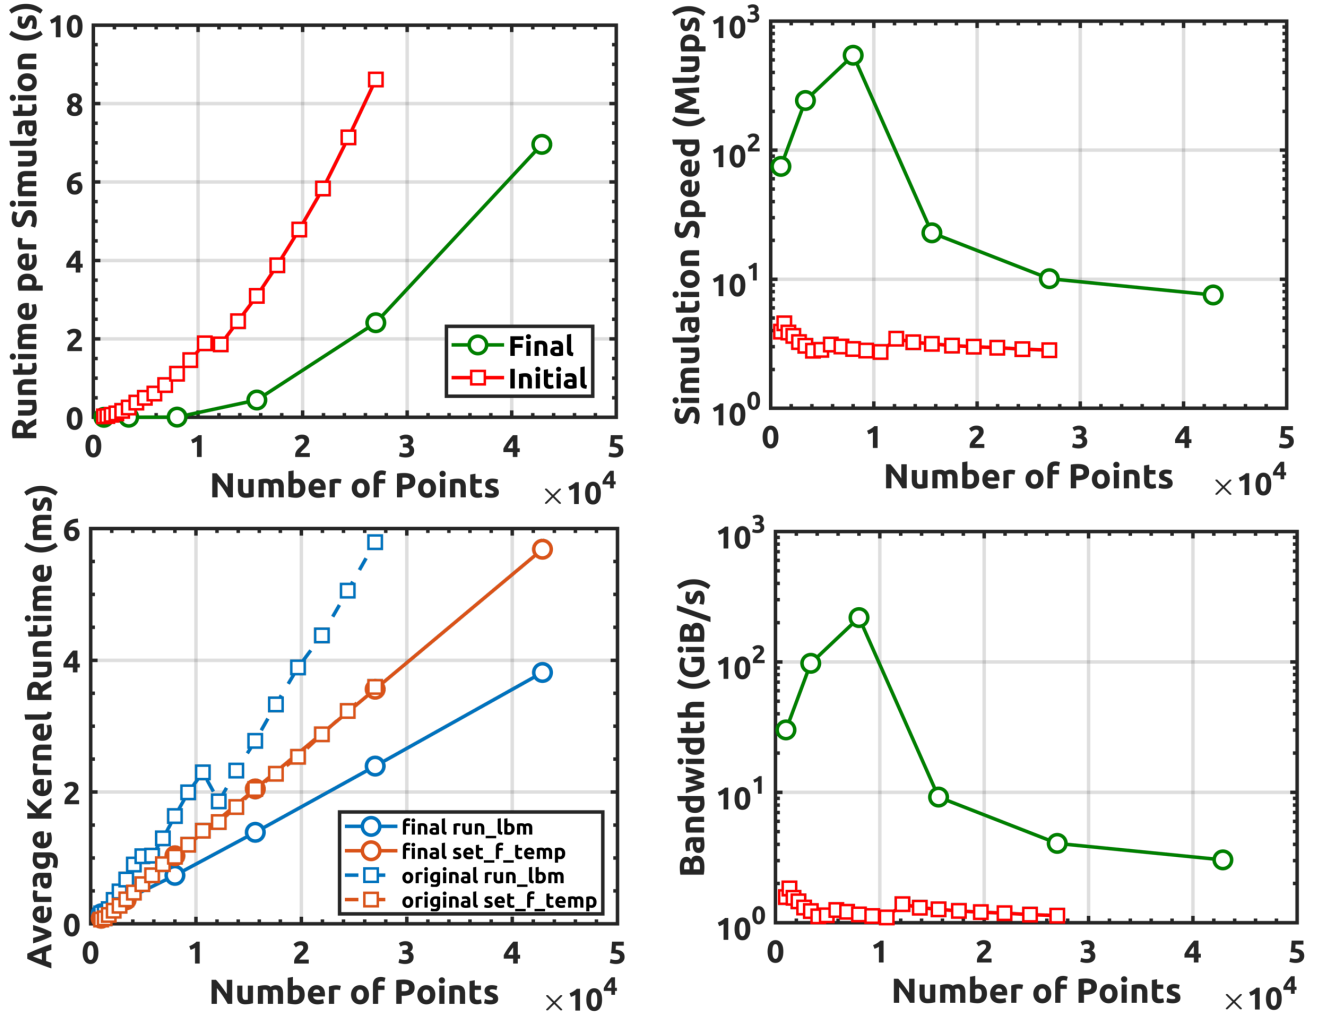


Figure 13: Comparison of important performance metrics between best-performing and original code.
result

Comparison with the prior submission shows the dramatic improvement in performance our new submission achieves. If n is the total number of points, then the total time for the final submission scales quadratically with problem size but at a slower rate than the initial submission. Similar to before, time to transfer in and out of device still scales in $O(n)$. The `set_ftemp` functions could be used as proxy for memory latency because it only has memory operations and as shown in the final figure, it scales in $O(n)$. It hasn't changed at all because the optimizations to remove the need for `set_ftemp` were unsuccessful. On the other hand, the final optimizations dramatically improved the performance of the `run_lbm` kernel. Optimization with grid and block sizes allowed dramatic increases in bandwidth because the sizes were optimized so as to allow more blocks to be run at a given time. All these optimizations dramatically increased the speed of the final submission especially because they combat the memory transfer and latency bottlenecks.

8.2 Analysis and Discussion of Improvements and Remaining Bottlenecks

The three main optimizations that were included in the final code was memory prefetching, optimizing the grid/block size, and removing the `CudaDeviceSynchronize`.

Fundamentally, the LBM method doesn't require any calls from CPU to device calls, and so the optimizations looked into were to minimize device memory read/writes and latency as well as grid size optimizations. Prefetching the f distributions

reduced the amount of memory reads, and the optimal grid sizing ensured there was enough shared memory for all the threads in a thread block so that storing in global memory could be avoided.

After further analysis, the compute operations were deemed to be already optimal and no further optimizations in this area could be thought of. Even in the original implementation that was submitted for check point 3, optimizations such as storing the output of divisions were precomputed and stored as much as possible. Thus, no additional compute optimizations were found.

As discussed in checkpoint 2, it was deemed that the primary avenue for optimizations with the Lattice Boltzmann Method should be to optimize the memory latency and transfers. The memory latency and transfer bottlenecks were as expected and optimizations in this area dramatically sped up performance. The analysis with varying grid sizes at the beginning of this report further confirmed this hunch. Additionally, memory transfer optimizations such as prefetching gave the expected boost in performance achieving a 20% in improvement. Similarly, optimizing the grid size ensured enough memory cache locally for threads so as to speed up performance.

8.3 Further Optimization

As mentioned before, if the push-pull implementation could have been debugged this would have been included in the final submission because this would have dramatically reduced the read/writes per iteration, which is especially beneficial as we have determined our algorithm is memory transfer and latency bound. A shared memory structure combined with grid size optimization could have been implemented so as to ensure that each thread block has enough memory to store distributions in shared memory by limiting the number of threads per block. But as is, the optimizations implemented in the final submission have lead to a dramatic speedup in speed compared to the original.

Further improvements that could be attempted to increase memory bandwidth include reorganizing the array of velocity densities and the kernels such that streaming could be performed across threads in a block. This would require the use of halo points which make this a complex endeavor, but the potential performance benefits from having data transfer at the L1 level might be worth the extra programming effort.

9. CONCLUSIONS AND SUMMARY

In this project, a serial implementation of the Lattice Boltzmann Method (LBM) was developed (see "serial" folder in [d8a7a7](#)). The Taylor-Green Vortex, an idealized flow whereby analytical solutions to the Navier-Stokes equations could be found, was used for validation of the implementation's correctness.

A parallel implementation was then performed, with the platform of choice being a GPU-based solution due to the high memory transfer requirements in the streaming operations involved in the LBM algorithm. The GPU implementation was optimized by:

1. Changing the grid-block structure to allow for more threads/warps to be executed at a given time, i.e. higher concurrency. This allowed for increased bandwidth and therefore reduced runtime (See Section 5).
2. Using local copies of frequently used variables within the CUDA kernels in order to reduce memory accesses to and from global *device* memory (See Section 4).
3. Removing the `cudaDeviceSynchronize` function with each iteration reduced the overall time of the simulation without affecting accuracy. This resulted in the largest performance improvement in this brief optimization of the LBM code. (See Section 6)