

Analysis Report

run_lbm(double[27]*, double[27]*, unsigned long, double, IndexBlock)

Duration	183.365 μ s
Grid Size	[11,1,1]
Block Size	[11,11,1]
Registers/Thread	60
Shared Memory/Block	0 B
Shared Memory Requested	48 KiB
Shared Memory Executed	48 KiB
Shared Memory Bank Size	4 B

[0] GeForce GTX 765M


GPU UUID	GPU-fe6c21ea-b7dc-9c6c-991c-33e89452d625
Compute Capability	3.0
Max. Threads per Block	1024
Max. Threads per Multiprocessor	2048
Max. Shared Memory per Block	48 KiB
Max. Shared Memory per Multiprocessor	48 KiB
Max. Registers per Block	65536
Max. Registers per Multiprocessor	65536
Max. Grid Dimensions	[2147483647, 65535, 65535]
Max. Block Dimensions	[1024, 1024, 64]
Max. Warps per Multiprocessor	64
Max. Blocks per Multiprocessor	16
Single Precision FLOP/s	1.325 TeraFLOP/s
Double Precision FLOP/s	55.2 GigaFLOP/s
Number of Multiprocessors	4
Multiprocessor Clock Rate	862.5 MHz
Concurrent Kernel	true
Max IPC	7
Threads per Warp	32
Global Memory Bandwidth	64.128 GB/s
Global Memory Size	1.952 GiB
Constant Memory Size	64 KiB
L2 Cache Size	256 KiB
Memcpy Engines	1
PCIe Generation	2
PCIe Link Rate	5 Gbit/s
PCIe Link Width	16

1. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results below indicate that the performance of kernel "run_lbm" is most likely limited by instruction and memory latency. You should first examine the information in the "Instruction And Memory Latency" section to determine how it is limiting performance.

1.1. Kernel Performance Is Bound By Instruction And Memory Latency

This kernel exhibits low compute throughput and memory bandwidth utilization relative to the peak performance of "GeForce GTX 765M". These utilization levels indicate that the performance of the kernel is most likely limited by the latency of arithmetic or memory operations. Achieved compute throughput and/or memory bandwidth below 60% of peak typically indicates latency issues.



2. Instruction and Memory Latency

Instruction and memory latency limit the performance of a kernel when the GPU does not have enough work to keep busy. The results below indicate that the GPU does not have enough work because the kernel does not execute enough blocks.

2.1. Grid Size Too Small To Hide Compute And Memory Latency

The kernel does not execute enough blocks to hide memory and operation latency. Typically the kernel grid size must be large enough to fill the GPU with multiple "waves" of blocks. Based on theoretical occupancy, device "GeForce GTX 765M" can simultaneously execute 8 blocks on each of the 4 SMs, so the kernel may need to execute a multiple of 32 blocks to hide the compute and memory latency. If the kernel is executing concurrently with other kernels then fewer blocks will be required because the kernel is sharing the SMs with those kernels.

Optimization: Increase the number of blocks executed by the kernel.

2.2. GPU Utilization May Be Limited By Register Usage

Theoretical occupancy is less than 100% but is large enough that increasing occupancy may not improve performance. You can attempt the following optimization to increase the number of warps on each SM but it may not lead to increased performance.


The kernel uses 60 registers for each thread (7260 registers for each block). This register usage is likely preventing the kernel from fully utilizing the GPU. Device "GeForce GTX 765M" provides up to 65536 registers for each block. Because the kernel uses 7260 registers for each block each SM is limited to simultaneously executing 8 blocks (32 warps). Chart "Varying Register Count" below shows how changing register usage will change the number of blocks that can execute on each SM.

Optimization: Use the `-maxrregcount` flag or the `__launch_bounds__` qualifier to decrease the number of registers used by each thread. This will increase the number of blocks that can execute on each SM. On devices with Compute Capability 5.2 turning global cache off can increase the occupancy limited by register usage.


Variable	Achieved	Theoretical	Device Limit	Grid Size: [11,1,1] (11 blocks) Block Size: [11,11,1] (121 threads)
Occupancy Per SM				
Active Blocks		8	16	
Active Warps	10.4	32	64	
Active Threads		1024	2048	
Occupancy	16.2%	50%	100%	
Warps				
Threads/Block		121	1024	
Warps/Block		4	32	
Block Limit		16	16	
Registers				
Registers/Thread		60	65536	
Registers/Block		8192	65536	
Block Limit		8	16	
Shared Memory				
Shared Memory/Block		0	49152	
Block Limit		0	16	

2.3. Occupancy Charts


The following charts show how varying different components of the kernel will impact theoretical occupancy.



Varying Register Count



Varying Shared Memory Usage



3. Compute Resources

GPU compute resources limit the performance of a kernel when those resources are insufficient or poorly utilized.

3.1. Function Unit Utilization


Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.

Load/Store - Load and store instructions for local, shared, global, constant, etc. memory.

Arithmetic - All arithmetic instructions including integer and floating-point add and multiply, logical and binary operations, etc.


Control-Flow - Direct and indirect branches, jumps, and calls.

Texture - Texture operations.




3.2. Instruction Execution Counts

The following chart shows the mix of instructions executed by the kernel. The instructions are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing instructions in that class. The "Inactive" result shows the thread executions that did not execute any instruction because the thread was predicated or inactive due to divergence.



3.3. Floating-Point Operation Counts

The following chart shows the mix of floating-point operations executed by the kernel. The operations are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing operations in that class. The results do not sum to 100% because non-floating-point operations executed by the kernel are not shown in this chart.

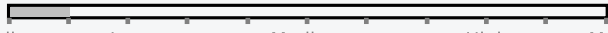
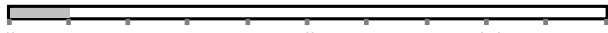





4. Memory Bandwidth

Memory bandwidth limits the performance of a kernel when one or more memories in the GPU cannot provide data at the rate requested by the kernel.

4.1. Memory Bandwidth And Utilization

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory.

Transactions	Bandwidth	Utilization	
L1/Shared Memory			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Global Loads	71874	13.512 GB/s	
Global Stores	71874	13.512 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	143748	27.025 GB/s	
L2 Cache			
L1 Reads	71874	13.512 GB/s	
L1 Writes	71874	13.512 GB/s	
Texture Reads	0	0 B/s	
Atomic	0	0 B/s	
Total	143748	27.025 GB/s	
Texture Cache			
Reads	0	0 B/s	
Device Memory			
Reads	40944	7.698 GB/s	
Writes	66788	12.556 GB/s	
Total	107732	20.254 GB/s	
System Memory			
[PCIe configuration: Gen2 x16, 5 Gbit/s]			
Reads	0	0 B/s	
Writes	2	376.003 kB/s	